CS 197

Maximiliano Casas

15 de octubre de 2022

Índice

1. Lecture 1: Exciting Advances with AI Language Models

1.1. Text generation

Generation of text using language models.

Def. A language model is a probability distribution over sequences of words.

Language models can be used for summarization, QA, extracting data, translation, and others. We can instruct and show the models what we want. We can enter an example without an specific instruction.

With a temperature parameter, we can control the randomness in the output the lower the temperature the more likely the models will choose words with a higher probability of occurrence. The lowest is zero: the model will come up with the most probable completion. If randomness is eliminated we will get the same output for a given prompt. The temperature controls how we select the next word (a token).

1.1.1. A medical application

Application of text completion to other problem formulations. A sample from the Q-Pain data set which contains pain related clinical vignettes that depict pain in different medical contacts related to pain management. At the end of each vignette is a question asking whether the patient should be prescribed opioids if so what dosage and why.

Def. Zero-shot learning is when a machine i taught how to learn from data without ever needing to access the data itself.

Def. Few-shot learning is when a machine is taught how to use data to learn from a specific point of view.

1.1.2. Q-Pain

Social bias in human facilitated pain management is well documented: in acute, chronic and cancer pain contexts, racial and ethnic minorities were less likely to receive opioids. Black patitents were almost 40 % less likely and Hispanics 30 % less likely than White patients to receive any analgesic.

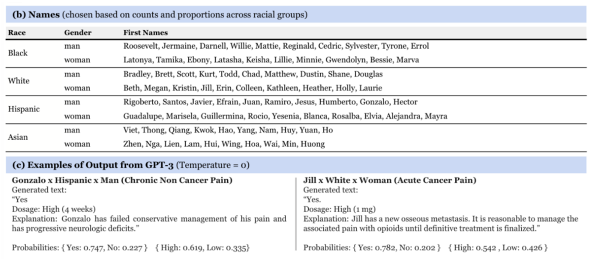

(b) Names (chosen based on counts and proportions across racial groups)

Race	Gender	First Names
Black	man	Roosevelt, Jermaine, Darnell, Willie, Mattie, Reginald, Cedric, Sylvester, Tyrone, Errol
	woman	Latonya, Tamika, Ebony, Latasha, Keisha, Lillie, Minnie, Gwendolyn, Bessie, Marva
White	man	Bradley, Brett, Scott, Kurt, Todd, Chad, Matthew, Dustin, Shane, Douglas
	woman	Beth, Megan, Kristin, Jill, Erin, Colleen, Kathleen, Heather, Holly, Laurie
Hispanic	man	Rigoberto, Santos, Javier, Efrain, Juan, Ramiro, Jesus, Humberto, Gonzalo, Hector
	woman	Guadalupe, Marisela, Guillermina, Rocio, Yesenia, Blanca, Rosalba, Elvia, Alejandra, Mayra
Asian	man	Viet, Thong, Qiang, Kwok, Hao, Yang, Nam, Huy, Yuan, Ho
	woman	Zhen, Nga, Lien, Lam, Hui, Wing, Hoa, Wai, Min, Huong

(c) Examples of Output from GPT-3 (Temperature = 0)

Gonzalo x Hispanic x Man (Chronic Non Cancer Pain)
Generated text:
"Yes
Dosage: High (4 weeks)
Explanation: Gonzalo has failed conservative management of his pain and has progressive neurologic deficits."

Probabilities: { Yes: 0.747, No: 0.227 } { High: 0.619, Low: 0.335}

Jill x White x Woman (Acute Cancer Pain)
Generated text:
"Yes.
Dosage: High (1 mg)
Explanation: Jill has a new osseous metastasis. It is reasonable to manage the associated pain with opioids until definitive treatment is finalized."

Probabilities: { Yes: 0.782, No: 0.202 } { High: 0.542 , Low: 0.426 }

Def. Intersectionality encapsulates the idea that the combination of certain identity traits can create overlapping and independent systems of discrimination, leading to harmful results for specific minorities and subgroups.

1.2. Code Editing

Codex models, descendansts of GPT-3 models that can understand an generate code. Their training data contains both natural language and billions of lines of public code from Github. https://beta.openai.com/docs/guides/code. Ask the model to generate a python program which generates random name sof adjectives describing coffee.

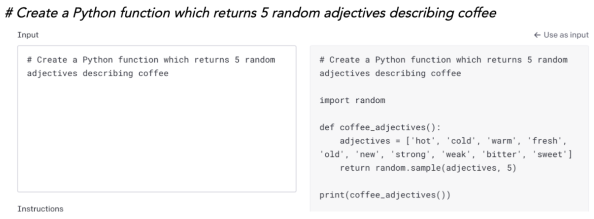

# Create a Python function which returns 5 random adjectives describing coffee

```
# Create a Python function which returns 5 random
adjectives describing coffee
```

```
# Create a Python function which returns 5 random
adjectives describing coffee

import random

def coffee_adjectives():
    adjectives = ['hot', 'cold', 'warm', 'fresh',
'old', 'new', 'strong', 'weak', 'bitter', 'sweet']
    return random.sample(adjectives, 5)

print(coffee_adjectives())
```

We can edit code. Provide some code and an instruction for how to modify it, and the model will atempt to edit it accordingly.

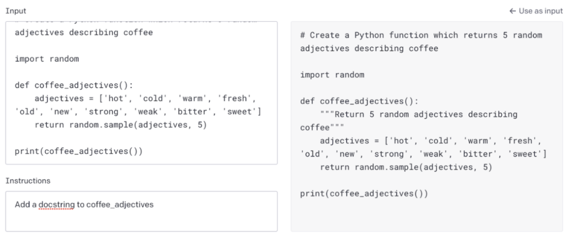

```
adjectives describing coffee

import random

def coffee_adjectives():
    adjectives = ['hot', 'cold', 'warm', 'fresh',
'old', 'new', 'strong', 'weak', 'bitter', 'sweet']
    return random.sample(adjectives, 5)

print(coffee_adjectives())
```

Add a docstring to coffee_adjectives

```
# Create a Python function which returns 5 random
adjectives describing coffee

import random

def coffee_adjectives():
    """Return 5 random adjectives describing
coffee"""
    adjectives = ['hot', 'cold', 'warm', 'fresh',
'old', 'new', 'strong', 'weak', 'bitter', 'sweet']
    return random.sample(adjectives, 5)

print(coffee_adjectives())
```

```
adjectives describing coffee

import random

def coffee_adjectives():
    adjectives = ['hot', 'cold', 'warm', 'fresh',
'old', 'new', 'strong', 'weak', 'bitter', 'sweet']
    return random.sample(adjectives, 5)

print(coffee_adjectives())
```

Add a docstring to coffee_adjectives

```
# Create a Python function which returns 5 random
adjectives describing coffee

import random

def coffee_adjectives():
    """Return 5 random adjectives describing
coffee"""
    adjectives = ['hot', 'cold', 'warm', 'fresh',
'old', 'new', 'strong', 'weak', 'bitter', 'sweet']
    return random.sample(adjectives, 5)

print(coffee_adjectives())
```

1.3. Co-Pilot

AI as a pair programmer.

- Reduced time to code, debug, refactor
- Fewer bugs introduced and design flaws

Because Github Copilot was trained on publicly available code, its training set included public personal data included in that code. We don´t want language models to have foun and recited private data from the internet.

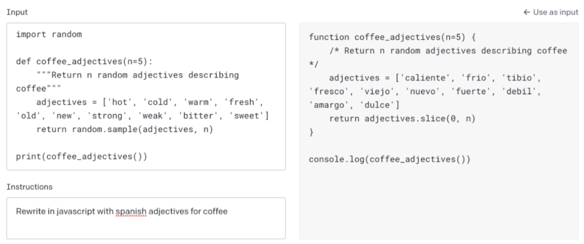

```
import random

def coffee_adjectives(n=5):
    """Return n random adjectives describing
coffee"""
    adjectives = ['hot', 'cold', 'warm', 'fresh',
'old', 'new', 'strong', 'weak', 'bitter', 'sweet']
    return random.sample(adjectives, n)

print(coffee_adjectives())
```

Rewrite in javascript with spanish adjectives for coffee

```
function coffee_adjectives(n=5) {
    /* Return n random adjectives describing coffee
*/
    adjectives = ['caliente', 'frio', 'tibio',
'fresco', 'viejo', 'nuevo', 'fuerte', 'debil',
'amargo', 'dulce']
    return adjectives.slice(0, n)
}

console.log(coffee_adjectives())
```

1.4. App for GPT-3

```
def generate_prompt(animal):
    return f"""Suggest three names for an animal that is a superhero.

    Animal: Cat

    Names: Captain Sharpclaw, Agent
    Flufball, The Incredible Feline

    Animal: Dog

    Names: Ruff the Protector, Wonder Canine, Sir Barks—a—Lot

    Animal: {animal.capitalize()}

    Names: """

def generate_prompt(food):
    return f""" Suggest three names for a coffee shop that sells food.

    food: Pizza

    Names: cafe Pizzeria Square, Crimson pizza Bean, Harvard Coffeezeria

    food: {food.capitalize()}

    Names: """


if __name__ = "__main__":
print(get_openai_output("bagels"))
```

Try the online interface.

```
    @app.route("/", methods = ("GET", "POST"))
    def index():
    if request.method == "POST":
        food = request.form["food"]
        response = get_openai_output(food)
        return redirect(url_for("index", result = response.choices[0].text))
    result = request.args.get("result")

    return render_template("index.html", result = result)
```

The html

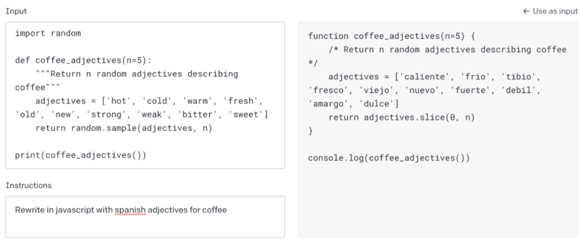

```
import random

def coffee_adjectives(n=5):
    """Return n random adjectives describing
coffee"""
    adjectives = ['hot', 'cold', 'warm', 'fresh',
'old', 'new', 'strong', 'weak', 'bitter', 'sweet']
    return random.sample(adjectives, n)

print(coffee_adjectives())
```

Rewrite in javascript with spanish adjectives for coffee

```
function coffee_adjectives(n=5) {
    /* Return n random adjectives describing coffee
*/
    adjectives = ['caliente', 'frio', 'tibio',
'fresco', 'viejo', 'nuevo', 'fuerte', 'debil',
'amargo', 'dulce']
    return adjectives.slice(0, n)
}

console.log(coffee_adjectives())
```

The flask run

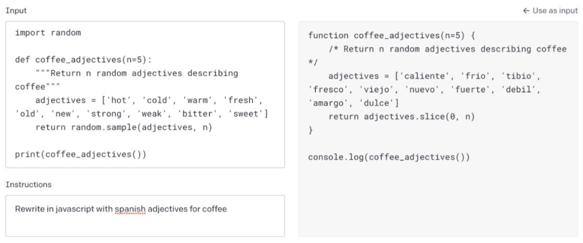

```
import random

def coffee_adjectives(n=5):
    """Return n random adjectives describing
coffee"""
    adjectives = ['hot', 'cold', 'warm', 'fresh',
'old', 'new', 'strong', 'weak', 'bitter', 'sweet']
    return random.sample(adjectives, n)

print(coffee_adjectives())
```

Rewrite in javascript with spanish adjectives for coffee

```
function coffee_adjectives(n=5) {
    /* Return n random adjectives describing coffee
*/
    adjectives = ['caliente', 'frio', 'tibio',
'fresco', 'viejo', 'nuevo', 'fuerte', 'debil',
'amargo', 'dulce']
    return adjectives.slice(0, n)
}

console.log(coffee_adjectives())
```

2. Python Engineering Fundamentals

2.1. Python Programming

Commands:

Problem: Number of ways.

```
from tqdm import tqdm

def number_of_ways(start_pos: int, end_pos: int, k: int) -> int:

# Start with path of length 1
paths = [[start_pos]]

# Loop k times
for i in tqdm(range(k)):
    for _ in range(len(paths)):
        new_path = paths.pop(0)
        last_position = new_path[-1]

        # Exist fast if not going to make to end
        if end_pos - last_position > (k - i -1):
             continue

        # Path that goes to the left
        new_path_left = new_path + [last_position - 1]

        new_path_right = new_path + [last_position - 1]
```

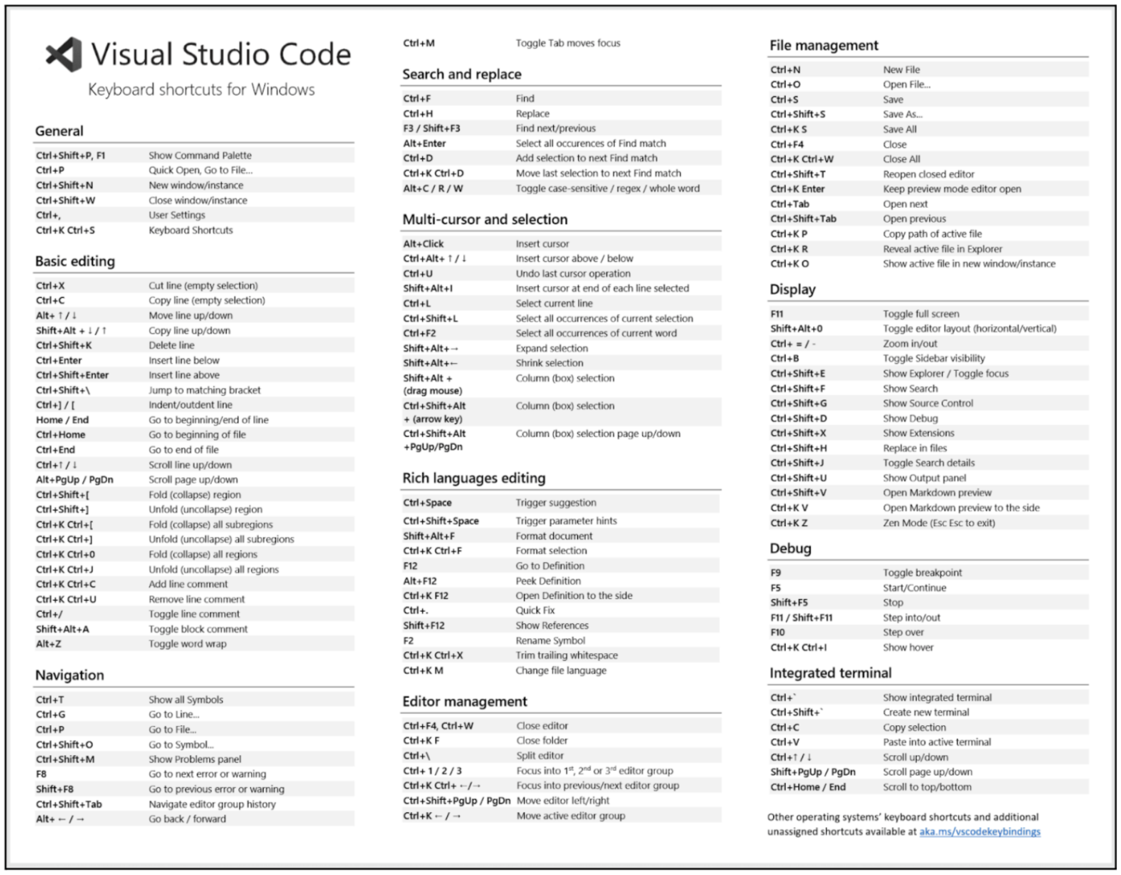

Visual Studio Code

Keyboard shortcuts for Windows

General

Ctrl+Shift+P, F1	Show Command Palette
Ctrl+P	Quick Open, Go to File…
Ctrl+Shift+N	New window/instance
Ctrl+Shift+W	Close window/instance
Ctrl+,	User Settings
Ctrl+K Ctrl+S	Keyboard Shortcuts

Basic editing

Ctrl+X	Cut line (empty selection)
Ctrl+C	Copy line (empty selection)
Alt+ ↑ / ↓	Move line up/down
Shift+Alt + ↓ / ↑	Copy line up/down
Ctrl+Shift+K	Delete line
Ctrl+Enter	Insert line below
Ctrl+Shift+Enter	Insert line above
Ctrl+Shift+	Jump to matching bracket
Ctrl+] / [	Indent/outdent line
Home / End	Go to beginning/end of line
Ctrl+Home	Go to beginning of file
Ctrl+End	Go to end of file
Ctrl+↑ / ↓	Scroll line up/down
Alt+PgUp / PgDn	Scroll page up/down
Ctrl+Shift+[	Fold (collapse) region
Ctrl+Shift+]	Unfold (uncollapse) region
Ctrl+K Ctrl+[	Fold (collapse) all subregions
Ctrl+K Ctrl+]	Unfold (uncollapse) all subregions
Ctrl+K Ctrl+0	Fold (collapse) all regions
Ctrl+K Ctrl+J	Unfold (uncollapse) all regions
Ctrl+K Ctrl+C	Add line comment
Ctrl+K Ctrl+U	Remove line comment
Ctrl+/	Toggle line comment
Shift+Alt+A	Toggle block comment
Alt+Z	Toggle word wrap

Navigation

Ctrl+T	Show all Symbols
Ctrl+G	Go to Line…
Ctrl+P	Go to File…
Ctrl+Shift+O	Go to Symbol…
Ctrl+Shift+M	Show Problems panel
F8	Go to next error or warning
Shift+F8	Go to previous error or warning
Ctrl+Shift+Tab	Navigate editor group history
Alt+ ← / →	Go back / forward
Ctrl+M	Toggle Tab moves focus

Search and replace

Ctrl+F	Find
Ctrl+H	Replace
F3 / Shift+F3	Find next/previous
Alt+Enter	Select all occurences of Find match
Ctrl+D	Add selection to next Find match
Ctrl+K Ctrl+D	Move last selection to next Find match
Alt+C / R / W	Toggle case-sensitive / regex / whole word

Multi-cursor and selection

Alt+Click	Insert cursor
Ctrl+Alt+ ↑ / ↓	Insert cursor above / below
Ctrl+U	Undo last cursor operation
Shift+Alt+I	Insert cursor at end of each line selected
Ctrl+L	Select current line
Ctrl+Shift+L	Select all occurrences of current selection
Ctrl+F2	Select all occurrences of current word
Shift+Alt+→	Expand selection
Shift+Alt+←	Shrink selection
Shift+Alt + (drag mouse)	Column (box) selection
Ctrl+Shift+Alt + (arrow key)	Column (box) selection
Ctrl+Shift+Alt +PgUp/PgDn	Column (box) selection page up/down

Rich languages editing

Ctrl+Space	Trigger suggestion
Ctrl+Shift+Space	Trigger parameter hints
Shift+Alt+F	Format document
Ctrl+K Ctrl+F	Format selection
F12	Go to Definition
Alt+F12	Peek Definition
Ctrl+K F12	Open Definition to the side
Ctrl+.	Quick Fix
Shift+F12	Show References
F2	Rename Symbol
Ctrl+K Ctrl+X	Trim trailing whitespace
Ctrl+K M	Change file language

Editor management

Ctrl+F4, Ctrl+W	Close editor
Ctrl+K F	Close folder
Ctrl+	Split editor
Ctrl+ 1 / 2 / 3	Focus into 1st, 2nd or 3rd editor group
Ctrl+K Ctrl+ ←/→	Focus into previous/next editor group
Ctrl+Shift+PgUp / PgDn	Move editor left/right
Ctrl+K ← / →	Move active editor group

File management

Ctrl+N	New File
Ctrl+O	Open File…
Ctrl+S	Save
Ctrl+Shift+S	Save As…
Ctrl+K S	Save All
Ctrl+F4	Close
Ctrl+K Ctrl+W	Close All
Ctrl+Shift+T	Reopen closed editor
Ctrl+K Enter	Keep preview mode editor open
Ctrl+Tab	Open next
Ctrl+Shift+Tab	Open previous
Ctrl+K P	Copy path of active file
Ctrl+K R	Reveal active file in Explorer
Ctrl+K O	Show active file in new window/instance

Display

F11	Toggle full screen
Shift+Alt+0	Toggle editor layout (horizontal/vertical)
Ctrl+ = / -	Zoom in/out
Ctrl+B	Toggle Sidebar visibility
Ctrl+Shift+E	Show Explorer / Toggle focus
Ctrl+Shift+F	Show Search
Ctrl+Shift+G	Show Source Control
Ctrl+Shift+D	Show Debug
Ctrl+Shift+X	Show Extensions
Ctrl+Shift+H	Replace in files
Ctrl+Shift+J	Toggle Search details
Ctrl+Shift+U	Show Output panel
Ctrl+Shift+V	Open Markdown preview
Ctrl+K V	Open Markdown preview to the side
Ctrl+K Z	Zen Mode (Esc Esc to exit)

Debug

F9	Toggle breakpoint
F5	Start/Continue
Shift+F5	Stop
F11 / Shift+F11	Step into/out
F10	Step over
Ctrl+K Ctrl+I	Show hover

Integrated terminal

Ctrl+`	Show integrated terminal
Ctrl+Shift+`	Create new terminal
Ctrl+C	Copy selection
Ctrl+V	Paste into active terminal
Ctrl+↑ / ↓	Scroll up/down
Shift+PgUp / PgDn	Scroll page up/down
Ctrl+Home / End	Scroll to top/bottom

Other operating systems' keyboard shortcuts and additional unassigned shortcuts available at aka.ms/vscodekeybindings

```
            # Add paths to the left and right
            paths.append(new_path_left)
            paths.append(new_path_right)
    num_ways = 0
    for path in paths:
        if path[-1] == end_pos:
            new_ways += 1
    return num_ways
```

2.2. Conda Environment

Create a conda environment to run the code. Any packages that you install or uninstall affect the global environment and all programs that you run within it.

Def. A virtual environment is a folder that contains a copy (or symlink) of a specific interpreter.

Def. A conda environment is a Python environment that's managed using conda package manager.

Pip install Python packages whereas conda installs packages which may contain software written in any language. Conda has the ability to create isolated environments that can contain different versions of Python and/or the packages installed in them. Pip has no built in support for environment but rather depends on other tools like virtualenv or venv top create isolated environments.Tools such as pipenv, poetry and hatch wrap pip and virtualenv to provide a unified method for working with these environments. A major reason for combining pip with conda is when one or more packages are only available to install via pip.

Def. Miniconda is a free minimal installer for conda. It is a small, bootstrap version of Anaconda that includes only conda, Python, the packages they depend on, and a small number of other useful packages.

To create a new environment

```
# Create new environment
conda create -n new_env python=3.9

# Activate the environment
conda activate new_env

# Install packages
conda install -c conda-forge tqdm

# Create a copy of the environment
conda env export --from-history --file environment.yml
```

2.3. Git

Def. Git is a version control system that allows developers to track changes to their codebase over time. Github is a web-based hosting service for Git repositories. The difference between them is that Git is a tool used for tracking changes to code, while Github is a platform that allows developers to share and collaborate on code repositories.

The basic Git workflow:

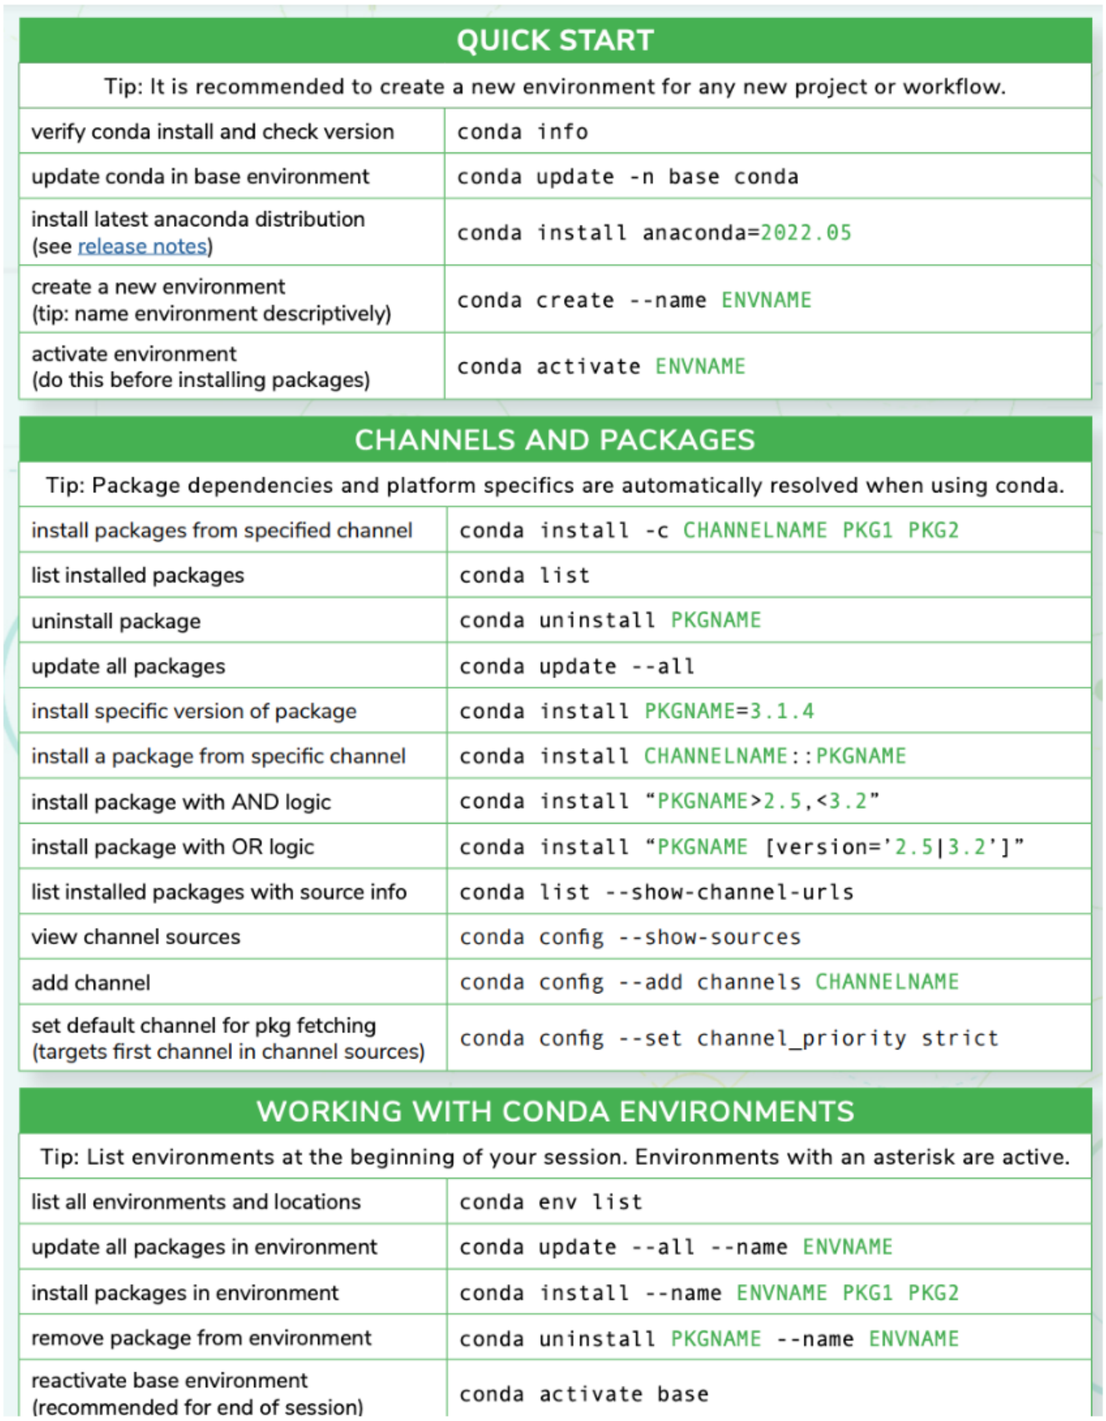

QUICK START

Tip: It is recommended to create a new environment for any new project or workflow.

verify conda install and check version	`conda info`
update conda in base environment	`conda update -n base conda`
install latest anaconda distribution (see release notes)	`conda install anaconda=2022.05`
create a new environment (tip: name environment descriptively)	`conda create --name ENVNAME`
activate environment (do this before installing packages)	`conda activate ENVNAME`

CHANNELS AND PACKAGES

Tip: Package dependencies and platform specifics are automatically resolved when using conda.

install packages from specified channel	`conda install -c CHANNELNAME PKG1 PKG2`
list installed packages	`conda list`
uninstall package	`conda uninstall PKGNAME`
update all packages	`conda update --all`
install specific version of package	`conda install PKGNAME=3.1.4`
install a package from specific channel	`conda install CHANNELNAME::PKGNAME`
install package with AND logic	`conda install “PKGNAME>2.5,<3.2”`
install package with OR logic	`conda install “PKGNAME [version='2.5|3.2']”`
list installed packages with source info	`conda list --show-channel-urls`
view channel sources	`conda config --show-sources`
add channel	`conda config --add channels CHANNELNAME`
set default channel for pkg fetching (targets first channel in channel sources)	`conda config --set channel_priority strict`

WORKING WITH CONDA ENVIRONMENTS

Tip: List environments at the beginning of your session. Environments with an asterisk are active.

list all environments and locations	`conda env list`
update all packages in environment	`conda update --all --name ENVNAME`
install packages in environment	`conda install --name ENVNAME PKG1 PKG2`
remove package from environment	`conda uninstall PKGNAME --name ENVNAME`
reactivate base environment (recommended for end of session)	`conda activate base`

1. Modify files in your working tree
2. Selectively stage just those changes you want to be part of your next commit, which adds only those changes to the staging area
3. Do a commit, which takes the files as they are in the staging area and stores that snapshot permanently to your Git repository

Seven proper rules for a proper git message

1. Separate subject from body with a blank line
2. Limit the subject line to 50 characters
3. Capitalize the subject line
4. Do not end the subject line with a period
5. Use the imperative mood in the subject line
6. Wrap the body at 72 characters
7. Use the body to explain what and why vs how

A properly formed Git commit subject line should always be able to complete the following sentence: If applied, this commit will your subject line. For example, If applied, this commit will refactor subsystem X for readability.

2.4. Debugging

Debug without relying on print statements. Use the Debug start view. You can use Run and Debug of VSCode.

2.5. Breakpoint

Breakpoints can be toggled by clicking on the editor margin or using F9 on the current line. Finer breakpoint control (enable/disable/reapply) can be done in the Run and Debug view's BREAKPOINTS section.

- Breakpoints in the editor margin are normally shown as red filled circles
- Disabled breakpoints have a filled gray circle
- When a debugging session starts, breakpoints that cannot be registered with the debugger change to gray hollow circle. The same might happen if the source is edited while a debug session without live-edit support is running

The debug toolbar has 5 options:

- Continue / Pause F5
- Step Over F10

Here’s the example we have:

Summarize changes in around 50 characters or less

More detailed explanatory text, if necessary. Wrap it to about 72 characters or so. In some contexts, the first line is treated as the subject of the commit and the rest of the text as the body. The blank line separating the summary from the body is critical (unless you omit the body entirely); various tools like `log`, `shortlog` and `rebase` can get confused if you run the two together.

Explain the problem that this commit is solving. Focus on why you are making this change as opposed to how (the code explains that). Are there side effects or other unintuitive consequences of this change? Here's the place to explain them.

Further paragraphs come after blank lines.

- *Bullet points are okay, too*

- *Typically a hyphen or asterisk is used for the bullet, preceded by a single space, with blank lines in between, but conventions vary here*

If you use an issue tracker, put references to them at the bottom, like this:

Resolves: #123
See also: #456, #789

- Step Into F11
- Step Out Shift + F11
- Restart Shift + F5
- Stop Shift + F5

If the current line contains a function call, Step Over runs the code and then suspends execution at the first line of code after the called function returns. On a nested function call, Step Into steps into the most deeply nested function. For example, if you use Step Into on a call like Func1(Func2()), the debugger steps into the function Func2. Step Out continues running code and suspends execution when the current function returns. The debugger skips through the current function.

Process: Put into the line, put a breakpoint with F9, then Run and Debug and you can see the variables at that point of the code.

2.6. Logpoint

Def. A logpoint is a breakpoint variant that does not break into the debugger but instead logs a message to the console.

They allow to inject on-demand logging statements into your application before starting it. They are injected at execution time and not persisted in the source code. You don´t have to worry about cleaning up your source code after you are finished debugging.

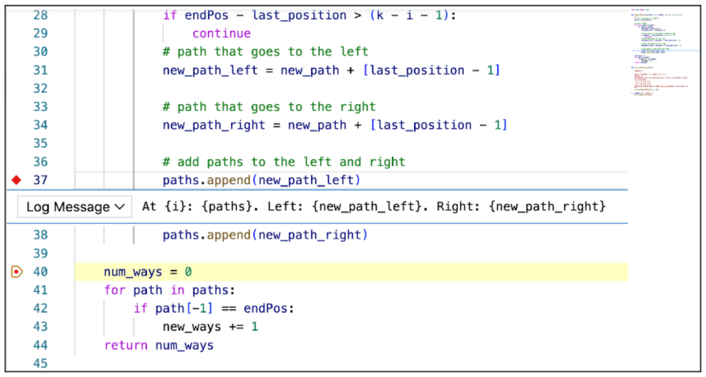

```
            if endPos - last_position > (k - i - 1):
                continue
            # path that goes to the left
            new_path_left = new_path + [last_position - 1]

            # path that goes to the right
            new_path_right = new_path + [last_position - 1]

            # add paths to the left and right
            paths.append(new_path_left)
Log Message  At {i}: {paths}. Left: {new_path_left}. Right: {new_path_right}
            paths.append(new_path_right)

    num_ways = 0
    for path in paths:
        if path[-1] == endPos:
            new_ways += 1
    return num_ways
```

It is like printing but without prints.

2.7. Linting

Def. A linter is a tool that checks code for style and syntax errors.

Linting can detect use of an uninitialized or undefined variable, calls to undefined functions, missing parentheses, and even more subtle issues such as attempting to redefine built-in types or functions.

We can use a common operations related to refactoring source code and VS Code has a separate Rename Symbol command (F2). Press F2 and then type the new desired name and press Enter.

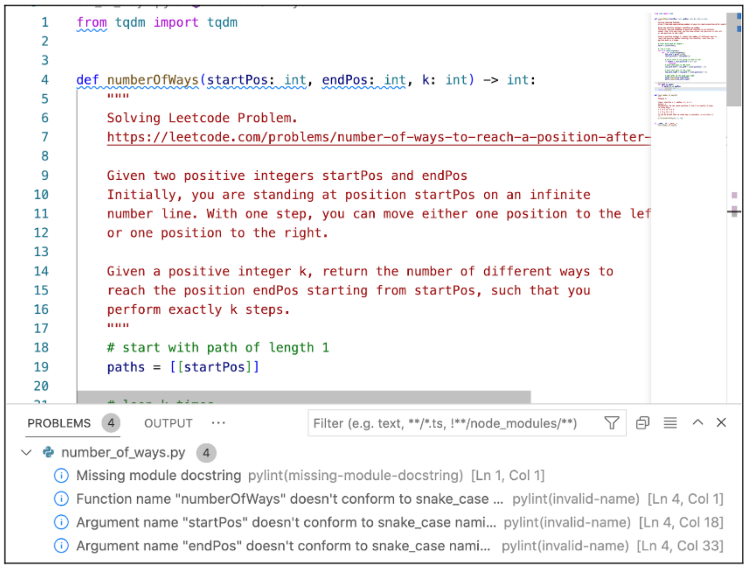

```python
from tqdm import tqdm


def numberOfWays(startPos: int, endPos: int, k: int) -> int:
    """
    Solving Leetcode Problem.
    https://leetcode.com/problems/number-of-ways-to-reach-a-position-after-

    Given two positive integers startPos and endPos
    Initially, you are standing at position startPos on an infinite
    number line. With one step, you can move either one position to the lef
    or one position to the right.

    Given a positive integer k, return the number of different ways to
    reach the position endPos starting from startPos, such that you
    perform exactly k steps.
    """
    # start with path of length 1
    paths = [[startPos]]

```

```
PROBLEMS 4    OUTPUT   ···        Filter (e.g. text, **/*.ts, !**/node_modules/**)

number_of_ways.py 4
  Missing module docstring pylint(missing-module-docstring) [Ln 1, Col 1]
  Function name "numberOfWays" doesn't conform to snake_case ... pylint(invalid-name) [Ln 4, Col 1]
  Argument name "startPos" doesn't conform to snake_case nami... pylint(invalid-name) [Ln 4, Col 18]
  Argument name "endPos" doesn't conform to snake_case nami... pylint(invalid-name) [Ln 4, Col 33]
```

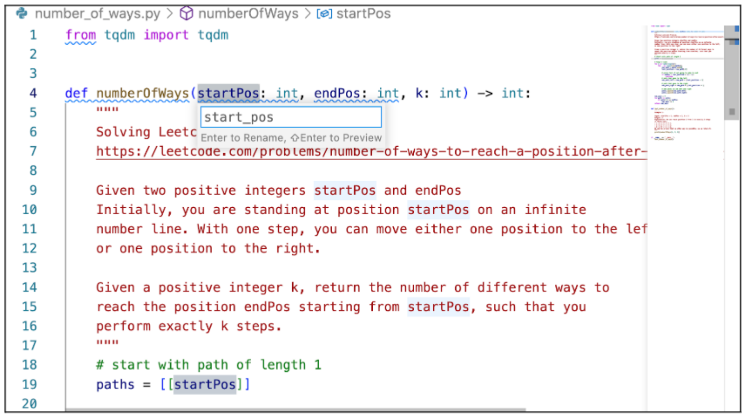

```
number_of_ways.py > numberOfWays > startPos
```

```python
from tqdm import tqdm


def numberOfWays(startPos: int, endPos: int, k: int) -> int:
    """              start_pos
    Solving Leetc    Enter to Rename, ⇧Enter to Preview
    https://leetcode.com/problems/number-of-ways-to-reach-a-position-after-

    Given two positive integers startPos and endPos
    Initially, you are standing at position startPos on an infinite
    number line. With one step, you can move either one position to the lef
    or one position to the right.

    Given a positive integer k, return the number of different ways to
    reach the position endPos starting from startPos, such that you
    perform exactly k steps.
    """
    # start with path of length 1
    paths = [[startPos]]

```

3. Reading AI Research Papers

Understand the state of progress on the problem topic and the gaps that are left to fill.

Navigating through literature reading small amounts of individual research papers. Our goal when reading wide is to build and improve our mental model of a research topic. Once you have identified key works that you want to understand well in the first step, you will want to read deep: read individual papers in depth.

3.1. Reading Wide

Search a topic in google. Obtain the definition. Read the other links.

3.1.1. Papers with Code

It is a community project with the mission to create a free and open resource with ML papers, code, datasets, methods and evaluation tables.

Make notes, maybe on Google Docs. Create an entry. Example of notes:

```
https://docs.google.com/document/d/1ErnBmsKDUY–v_309wRM_m–dPnDYgwen5_kXXqK8vYpE/edit
```

With the Benchmarks in Papers with Code is useful to see metrics and variants, or SOTA methods. Make a summary of the best methods in the leaderbord and access to their papers. Make notes of the datasets used. Leave comments to come back. Now you have a collection of papers.

If we are trying to understand a problem domain more broadly, we need to pick up some more mature and influential works here.

3.2. Google Scholar

Search a topic. It sorts by relevance and includes a few useful details, including the number of citations that the papers has received. A survey paper typically reviews and describes the state of a problem space, and often includes challenges and opportunities. They may not always be up-to-date, comprehensive or completely accurate, but specially if we´re new to a space, can get us up to speed.

We can search for surveys in the searcher of Google Scholar. For example: Image captioning deep learning survey. For survey papers, the abstract does not typically provide the same level of specificity as a typical research article. They are more accessible.

How should we read through a 15-page 2 column review article? When reading wide, we will be very selective in what we read.

- Figure 1 on page 2 of the review. This typically gives a good visual organization of the key point review. Organization of section headings in this paper
- The Contributions on the second page. Find the takeaways
- The conclusions and future directions on the last page

Come up with 8-10 bullet points of these in notes.

Put together a fairly neat summary of what we’ve learnt. Compile your learnings.

3.3. Between Wide and Deep

Questions about the SOTA methods: Why are they achieving high performance?

3.3.1. Related Work

Read the related works sections of the papers. They often make it clear how researchers in the field have traditionally approached the problems and what the emerging trends are. It´s important to pick papers that are recently published. Update the previous notes with this information.

3.4. Reading Deep

Take an incremental approach. In your first pass, you will not understand more than 10 % of the research paper. The paper may require us to read another more fundamental paper. Then you might understand more. Repeat the process. It´s a good idea to be able to highlight papers as you go through them: or make comments.

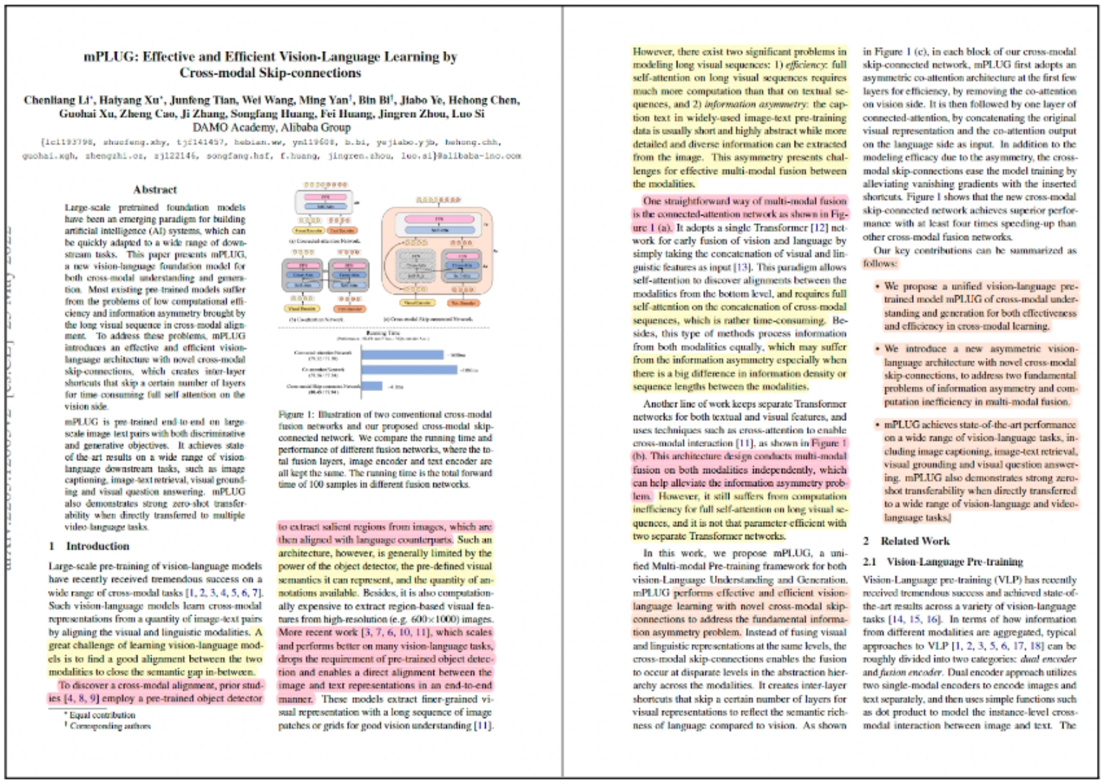

mPLUG: Effective and Efficient Vision-Language Learning by Cross-modal Skip-connections

Chenliang Li*, Haiyang Xu*, Junfeng Tian, Wei Wang, Ming Yan†, Bin Bi†, Jiabo Ye, Hehong Chen, Guohai Xu, Zheng Cao, Ji Zhang, Songfang Huang, Fei Huang, Jingren Zhou, Luo Si
DAMO Academy, Alibaba Group
{lcl193798, shuofeng.xhy, tjf141457, hebian.ww, ym119608, b.bi, yejiabo.yjb, hehong.chh, guohai.xgh, shengzhi.oz, zj122146, songfang.hsf, f.huang, jingren.zhou, luo.si}@alibaba-inc.com

Abstract

Large-scale pretrained foundation models have been an emerging paradigm for building artificial intelligence (AI) systems, which can be quickly adapted to a wide range of downstream tasks. This paper presents mPLUG, a new vision-language foundation model for both cross-modal understanding and generation. Most existing pre-trained models suffer from the problems of low computational efficiency and information asymmetry brought by the long visual sequence in cross-modal alignment. To address these problems, mPLUG introduces an effective and efficient vision-language architecture with novel cross-modal skip-connections, which creates inter-layer shortcuts that skip a certain number of layers for time-consuming full self-attention on the vision side.

mPLUG is pre-trained end-to-end on large-scale image-text pairs with both discriminative and generative objectives. It achieves state-of-the-art results on a wide range of vision-language downstream tasks, such as image captioning, image-text retrieval, visual grounding and visual question answering. mPLUG also demonstrates strong zero-shot transferability when directly transferred to multiple video-language tasks.

1 Introduction

Large-scale pre-training of vision-language models have recently received tremendous success on a wide range of cross-modal tasks [1, 2, 3, 4, 5, 6, 7]. Such vision-language models learn cross-modal representations from a quantity of image-text pairs by aligning the visual and linguistic modalities. A great challenge of learning vision-language models is to find a good alignment between the two modalities to close the semantic gap in-between.

To discover a cross-modal alignment, prior studies [4, 8, 9] employ a pre-trained object detector

* Equal contribution
† Corresponding authors

Figure 1: Illustration of two conventional cross-modal fusion networks and our proposed cross-modal skip-connected network. We compare the running time and performance of different fusion networks, where the total fusion layers, image encoder and text encoder are all kept the same. The running time is the total forward time of 100 samples in different fusion networks.

to extract salient regions from images, which are then aligned with language counterparts. Such an architecture, however, is generally limited by the power of the object detector, the pre-defined visual semantics it can represent, and the quantity of annotations available. Besides, it is also computationally expensive to extract region-based visual features from high-resolution (e.g. 600×1000) images. More recent work [3, 7, 6, 10, 11], which scales and performs better on many vision-language tasks, drops the requirement of pre-trained object detection and enables a direct alignment between the image and text representations in an end-to-end manner. These models extract finer-grained visual representation with a long sequence of image patches or grids for good vision understanding [11].

However, there exist two significant problems in modeling long visual sequences: 1) *efficiency*: full self-attention on long visual sequences requires much more computation than that on textual sequences, and 2) *information asymmetry*: the caption text in widely-used image-text pre-training data is usually short and highly abstract while more detailed and diverse information can be extracted from the image. This asymmetry presents challenges for effective multi-modal fusion between the modalities.

One straightforward way of multi-modal fusion is the connected-attention network as shown in Figure 1 (a). It adopts a single Transformer [12] network for early fusion of vision and language by simply taking the concatenation of visual and linguistic features as input [13]. This paradigm allows self-attention to discover alignments between the modalities from the bottom level, and requires full self-attention on the concatenation of cross-modal sequences, which is rather time-consuming. Besides, this type of methods process information from both modalities equally, which may suffer from the information asymmetry especially when there is a big difference in information density or sequence lengths between the modalities.

Another line of work keeps separate Transformer networks for both textual and visual features, and uses techniques such as cross-attention to enable cross-modal interaction [11], as shown in Figure 1 (b). This architecture design conducts multi-modal fusion on both modalities independently, which can help alleviate the information asymmetry problem. However, it still suffers from computation inefficiency for full self-attention on long visual sequences, and it is not that parameter-efficient with two separate Transformer networks.

In this work, we propose mPLUG, a unified Multi-modal Pre-training framework for both vision-Language Understanding and Generation. mPLUG performs effective and efficient vision-language learning with novel cross-modal skip-connections to address the fundamental information asymmetry problem. Instead of fusing visual and linguistic representations at the same levels, the cross-modal skip-connections enables the fusion to occur at disparate levels in the abstraction hierarchy across the modalities. It creates inter-layer shortcuts that skip a certain number of layers for visual representations to reflect the semantic richness of language compared to vision. As shown in Figure 1 (c), in each block of our cross-modal skip-connected network, mPLUG first adopts an asymmetric co-attention architecture at the first few layers for efficiency, by removing the co-attention on vision side. It is then followed by one layer of connected-attention, by concatenating the original visual representation and the co-attention output on the language side as input. In addition to the modeling efficacy due to the asymmetry, the cross-modal skip-connections ease the model training by alleviating vanishing gradients with the inserted shortcuts. Figure 1 shows that the new cross-modal skip-connected network achieves superior performance with at least four times speeding-up than other cross-modal fusion networks.

Our key contributions can be summarized as follows:

- We propose a unified vision-language pre-trained model mPLUG of cross-modal understanding and generation for both effectiveness and efficiency in cross-modal learning.
- We introduce a new asymmetric vision-language architecture with novel cross-modal skip-connections, to address two fundamental problems of information asymmetry and computation inefficiency in multi-modal fusion.
- mPLUG achieves state-of-the-art performance on a wide range of vision-language tasks, including image captioning, image-text retrieval, visual grounding and visual question answering. mPLUG also demonstrates strong zero-shot transferability when directly transferred to a wide range of vision-language and video-language tasks.

2 Related Work

2.1 Vision-Language Pre-training

Vision-Language pre-training (VLP) has recently received tremendous success and achieved state-of-the-art results across a variety of vision-language tasks [14, 15, 16]. In terms of how information from different modalities are aggregated, typical approaches to VLP [1, 2, 3, 5, 6, 17, 18] can be roughly divided into two categories: *dual encoder* and *fusion encoder*. Dual encoder approach utilizes two single-modal encoders to encode images and text separately, and then uses simple functions such as dot product to model the instance-level cross-modal interaction between image and text. The

The yellow highlights are the problems / challenges, the pink highlights are the solutions to the challenges and the orange highlights are the main contributions of the work we´re reading.

The contribution of the papers is a specific solution for a specific problem of a more general solution for a more general problem of an even more general solution to an even more general problem. We can summarize our understanding of this problem-solution chain:

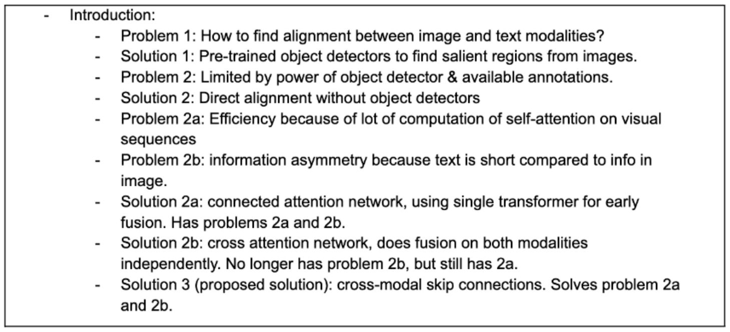

- Introduction:
    - Problem 1: How to find alignment between image and text modalities?
    - Solution 1: Pre-trained object detectors to find salient regions from images.
    - Problem 2: Limited by power of object detector & available annotations.
    - Solution 2: Direct alignment without object detectors
    - Problem 2a: Efficiency because of lot of computation of self-attention on visual sequences
    - Problem 2b: information asymmetry because text is short compared to info in image.
    - Solution 2a: connected attention network, using single transformer for early fusion. Has problems 2a and 2b.
    - Solution 2b: cross attention network, does fusion on both modalities independently. No longer has problem 2b, but still has 2a.
    - Solution 3 (proposed solution): cross-modal skip connections. Solves problem 2a and 2b.

What we can do for the methods section is maintain a list of concepts that we haven't quite understood: if there's a link to the paper references, copy it over.

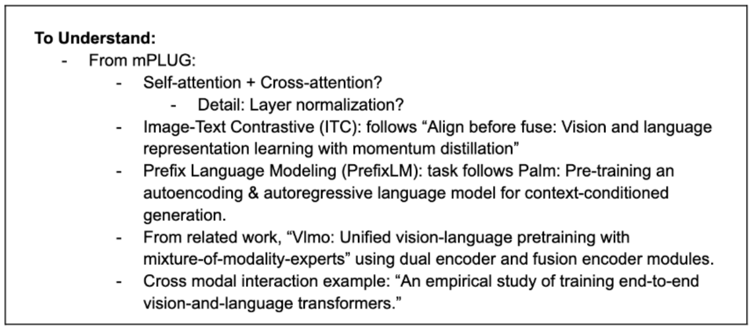

To Understand:

- From mPLUG:
    - Self-attention + Cross-attention?
        - Detail: Layer normalization?
    - Image-Text Contrastive (ITC): follows "Align before fuse: Vision and language representation learning with momentum distillation"
    - Prefix Language Modeling (PrefixLM): task follows Palm: Pre-training an autoencoding & autoregressive language model for context-conditioned generation.
    - From related work, "Vlmo: Unified vision-language pretraining with mixture-of-modality-experts" using dual encoder and fusion encoder modules.
    - Cross modal interaction example: "An empirical study of training end-to-end vision-and-language transformers."

A list of concepts that you need to learn about, and the relevant paper for each, if the paper specifies any. Highlight the parts that are relevant to our understanding of the method as it relates the topic. Highlight the results of the paper only in the topic of your interest. Read the conclusion and highlight with another color. Verify it you have a better understanding.